

**CS 211**

**Computer Architecture**

**Fall 2021**

David Menendez

# Multidimensional Arrays

```
int m[3][3];
```

**m** is an array of arrays of ints


**m[0]** is an array of ints

**m[0][0]** is an int

**sizeof(m) = ?**

**sizeof(m[0]) = ?**

**sizeof(m[0][0]) = ?**



# Multidimensional Arrays

- To access an element, the compiler needs to know the length of the rows
  - e.g., `int m[ROWS][COLS]`; the row length is COLS
  - `m[row][col]` means  $*(m + row * COLS + col)$ ;
  - `int v[X_RANGE][Y_RANGE][Z_RANGE]`;
  - `v[x][y][z]` means  $*(v + (x * Y\_RANGE + y) * Z\_RANGE + z)$ ;
- The compiler knows this because of the declaration

# Multidimensional Arrays

- How can we pass a multidimensional array to a function?
  - `void transpose(double matrix[][3]);`
  - `void something_else(int voxels[][100][100]);`
  - First dimension can be dropped because it isn't needed; others must be given

# Double Pointers

- Single pointers and single-dimension arrays can be mostly used interchangeably
- The differences come up in higher dimensions
  - `char text[X][Y]` is an array of arrays
    - contiguous in memory, rectangular
  - `char **argv` is a pointer to a pointer
    - No further requirements

# Double Pointers

- Allocating a 2D array with double pointers:


```
double **matrix = malloc(rows * sizeof(double *));
```

```
for (i = 0; i < rows; ++i) {
```

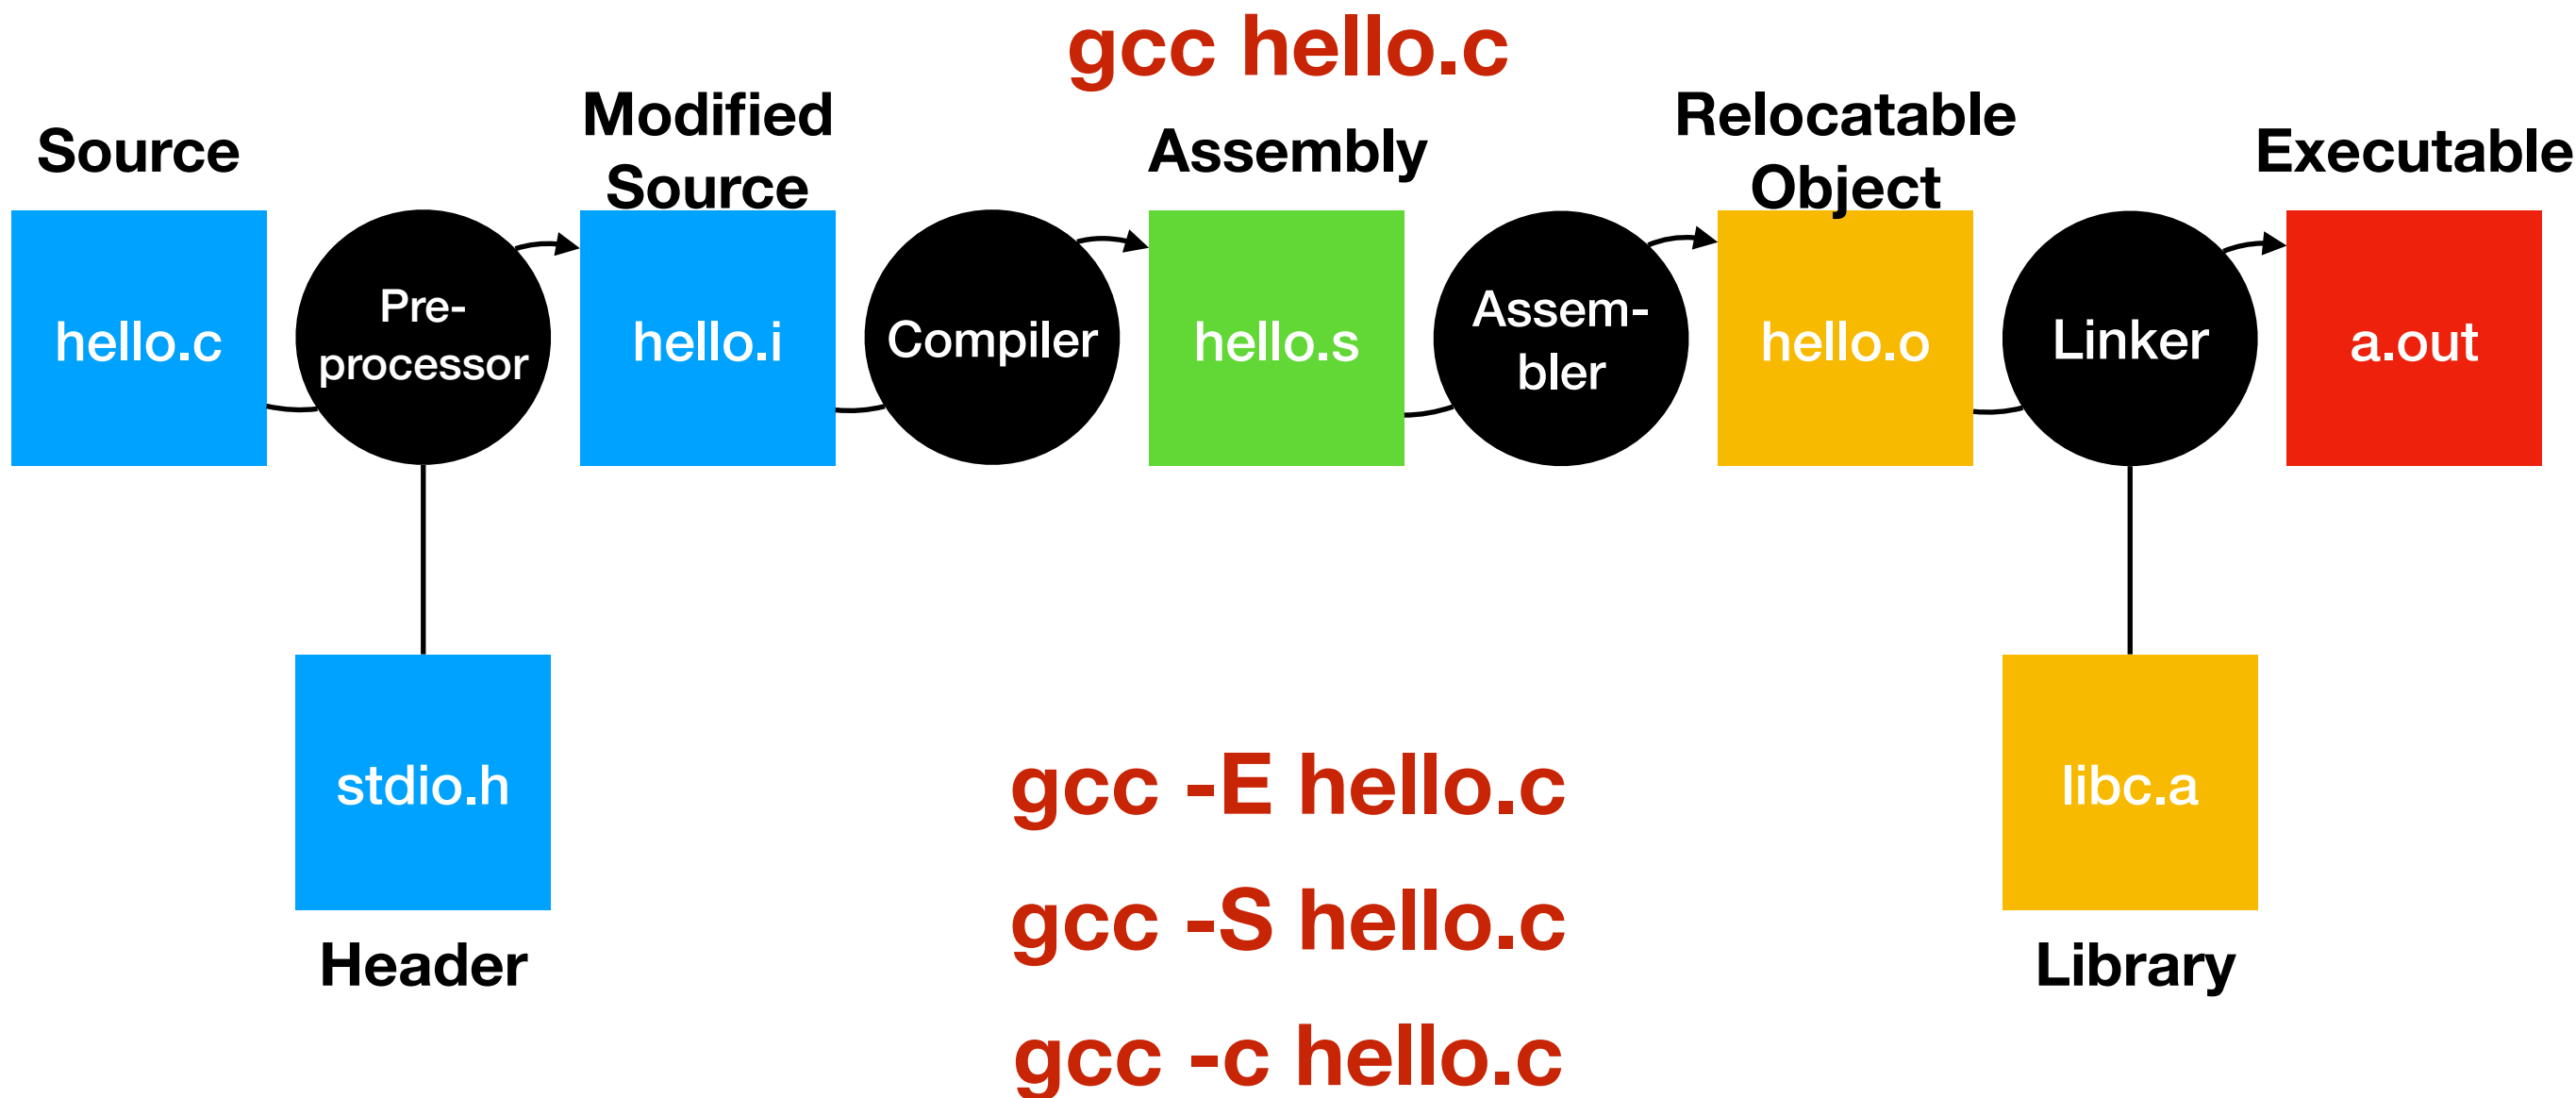
```
    matrix[i] = malloc(cols * sizeof(double));
```

```
}
```

# Compilation

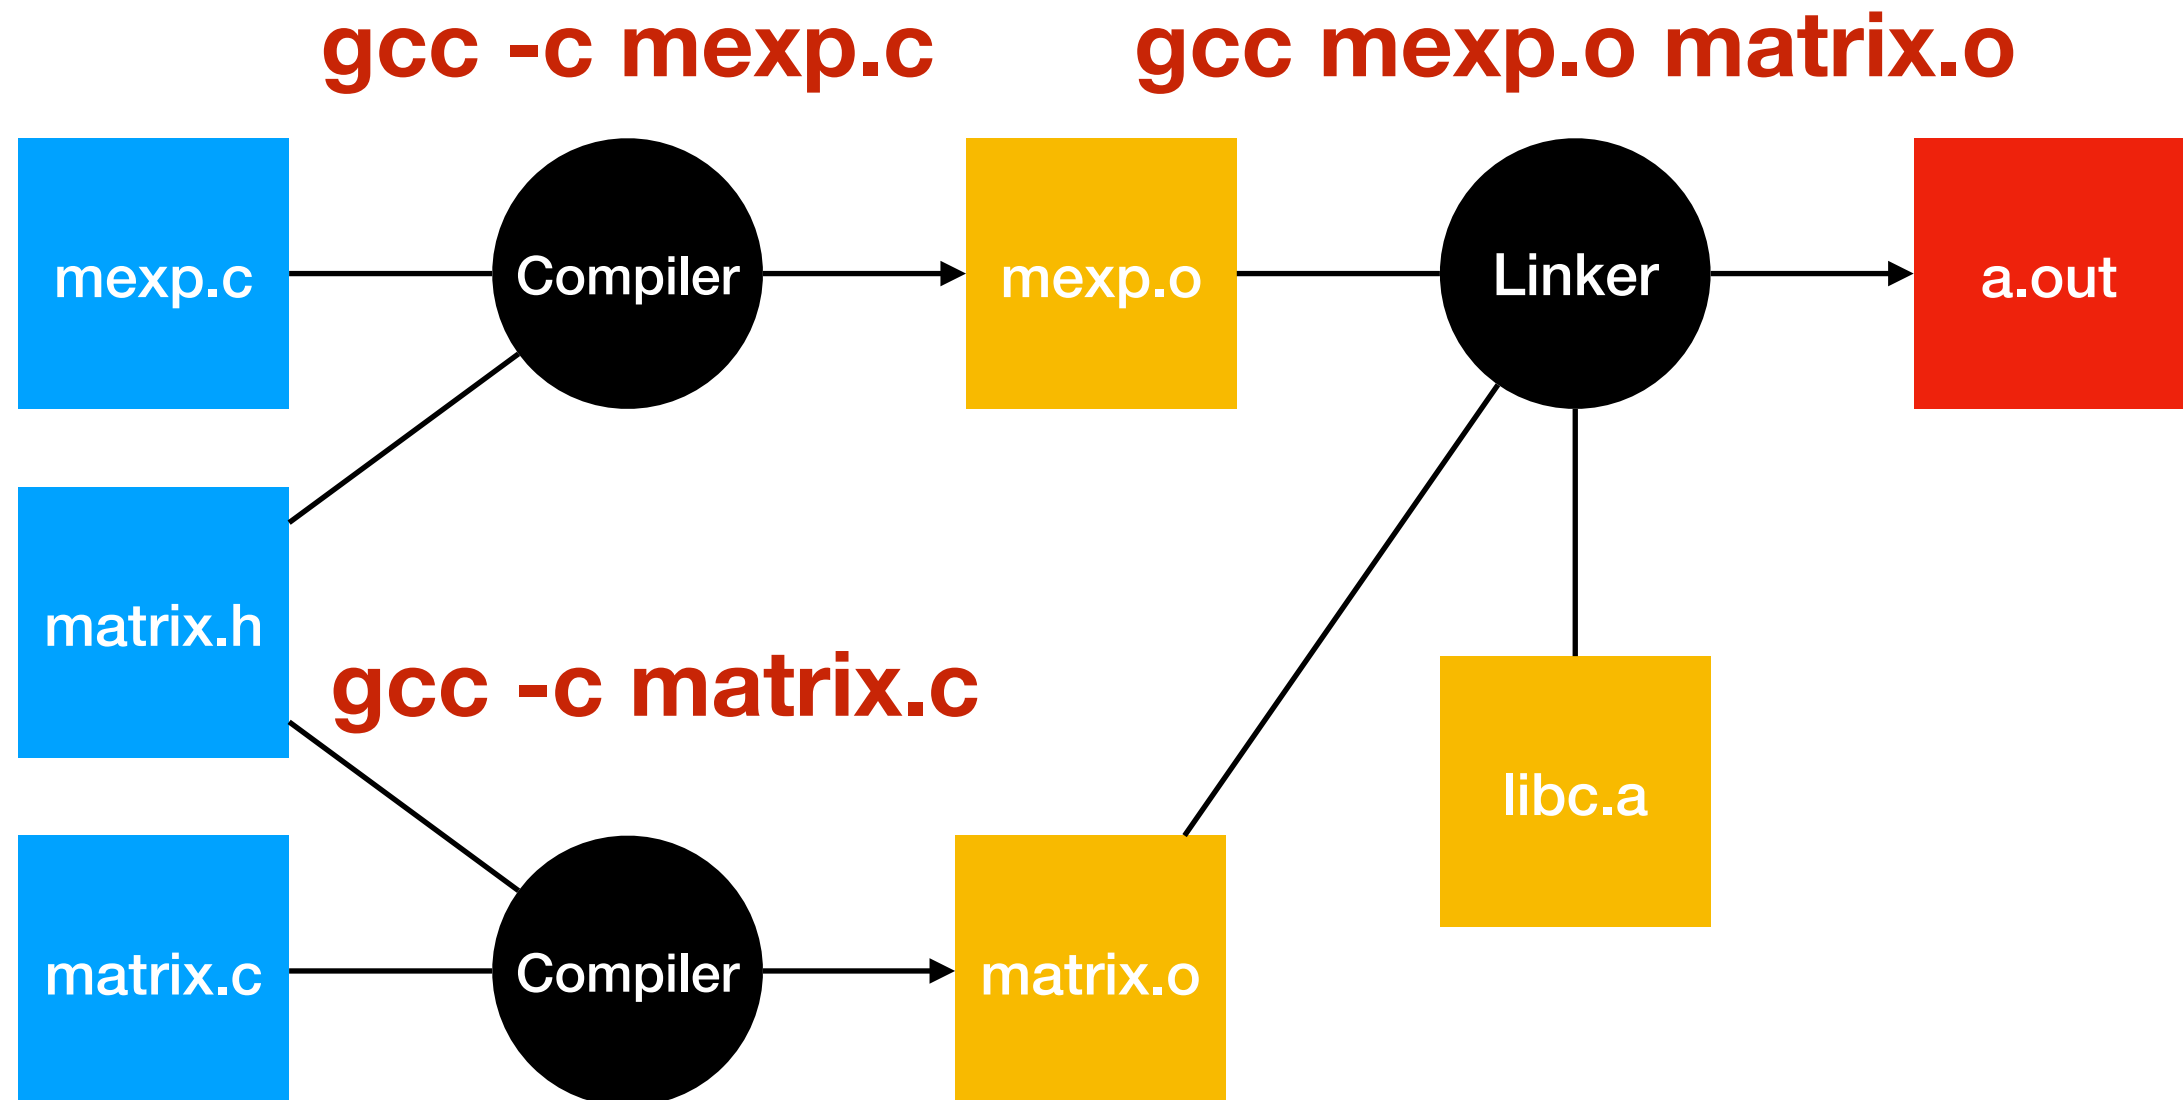


# Compilation





# Separate Compilation



# Running on Hardware

- What is the interface between software and hardware?
  - ISA: Instruction Set Architecture
- How do we represent data?
  - Data types (integers, floats, pointers, etc.)
  - Instructions themselves are data

# What Computers Do

- Computers manipulate data
  - Specifically numbers
  - Most things can be represented using numbers (words, images, sound, etc.)
- Numbers are abstract: we need a representation to manipulate them
  - Many notations used through history
  - Today, decimal notation (base 10) is most common

# Base 10: Decimal

- You already know this one
- The right-most (least significant) digit is the 1s place, to its left is the 10s place, to its left is the 100s place, etc.
  - $86042 = 80000 + 6000 + 000 + 40 + 2$
  - $86042 = 8 \times 10^4 + 6 \times 10^3 + 0 \times 10^2 + 4 \times 10^1 + 2 \times 10^0$
- We multiply each digit by the  $n$ th power of 10 (the base), where  $n$  is 0 for the right-most digit and increases for each digit to the left
- Side note: this means that initial zeros are irrelevant
  - $100 = 0100 = 00100 = \dots$

# Base $n$ Notation

- Generalization of decimal notation
- Numbers written as a sequence of digits
  - $d_k d_{k-1} \dots d_2 d_1 d_0$
- Each digit is multiplied by a place value
  - $x = d_k n^k + \dots + d_1 n^1 + d_0 n^0$
- Digit values typically range from 0 to  $n-1$

# Base 2: Binary

- In binary notation, the only digits are 0 and 1
  - bit  $\Leftarrow$  “binary digit”
- $x = d_k 2^k + \dots + d_1 2^1 + d_0 2^0$
- Binary is great for computers
  - Easy to represent with switches (on/off)
  - Can be manipulated by digital logic in hardware
- Binary is less good for humans
  - Relatively small numbers require a lot of digits

# Base 16: Hexadecimal

- Digits are 0–9 and A–F (representing 10–15)
- $x = d_k 16^k + \dots d_1 16^1 + d_0 16^0$
- More compact than binary
  - $16 = 2^4$ , so 1 hex digit is 4 bits
  - Bytes are two hex digits (00–FF)

# Hex-Binary Conversion

| Hex | Binary | Hex | Binary |
|-----|--------|-----|--------|
| 0   | 0000   | 8   | 1000   |
| 1   | 0001   | 9   | 1001   |
| 2   | 0010   | A   | 1010   |
| 3   | 0011   | B   | 1011   |
| 4   | 0100   | C   | 1100   |
| 5   | 0101   | D   | 1101   |
| 6   | 0110   | E   | 1110   |
| 7   | 0111   | F   | 1111   |



# Base 8: Octal

- Digits are 0–7
- $x = d_k 8^k + \dots d_0 8^0$
- More compact than binary
  - One octal digit is 3 bits
- Few things divide into groups of 3 bits these days

# Decimal–Binary Conversion

|             |             |
|-------------|-------------|
| <b>x=13</b> | <b>1</b>    |
| <b>x=6</b>  | <b>01</b>   |
| <b>x=3</b>  | <b>101</b>  |
| <b>x=1</b>  | <b>1101</b> |

- If  $x$  is odd, the last bit is 1, otherwise it is 0
- $x \leftarrow x \div 2$ . If  $x$  is odd, the next-to-last bit is 1, otherwise 0
- $x \leftarrow x \div 2$ . If  $x$  is odd, ...

# Other Conversions

- In general, repeatedly divide by  $n$  and take remainder
- Remainder gives digits  $0-(n-1)$  starting with  $d_0$
- How would you convert decimal to octal?
- How would you convert binary to base 12?

# Base $n$ Rationals

- The decimal point extends base 10 integers with additional places smaller than 1 (e.g., 1/10ths place, 1/100ths place)
- Similar “radix points” can be used in base  $n$
- Digits after the radix point are multiplied by negative powers of  $n$ 
  - $x = \dots + d_1n^1 + d_0n^0 + d_{-1}n^{-1} + d_{-2}n^{-2} + \dots$
- 101.11 in binary is  $4+0+1+\frac{1}{2}+\frac{1}{4} = 5.75$

# Value vs Notation

- Decimal, binary, hexadecimal, etc. are ways of *writing* numbers
- Computers use binary to encode integers
- All arithmetic operates on binary integers
- printf and others convert to decimal etc. when printing
  - Hence %d, %x, %o, etc.
- There are no “hex ints”—the value of any integer can be written in any convenient notation

# Data Sizes

- Primitive number types use a fixed number of digits
  - Usually multiples of 8 (8 bits = 1 byte)
- Types vary in C, but char = 1 byte, short int  $\geq 2$  bytes, long int  $\geq 4$  bytes

| Type      | 16-bit | 32-bit | 64-bit |
|-----------|--------|--------|--------|
| char      | 1      | 1      | 1      |
| short int | 2      | 2      | 2      |
| int       | 2      | 4      | 4      |
| long int  | 4      | 4      | 8      |
| void *    | 2      | 4      | 8      |
| float     | 4      | 4      | 4      |
| double    | 8      | 8      | 8      |

**Typical, but  
not universal**

# Big- and Little-Endian

- Large types consist of multiple bytes
- A 4-byte integer at address  $A$  will use bytes  $A \dots A+3$ 
  - But what will be in the specific bytes?
- Consider a large number  $AB10CD2F_{16}$
- Big Endian: most significant byte at smallest address:
  - AB 10 CD 2F
- Little Endian: least significant byte at smallest address:
  - 2F CD 10 AB

# When Does Endianness Matter?

- The CPU is designed to work with whichever system it has
  - E.g., reads, writes, arithmetic will all work correctly
- But what about sending data to other computers?
  - Need a standard representation for compatibility
  - Most network standards are big-endian
  - Some file formats have byte-order marks



# Encoding Data

- How could we encode a playing card?
  - 4 suits, 13 ranks = 52 cards
  - Operation: compare two cards of the same suit
- One byte each for suit and rank
- One byte: 4 bits for rank, 2 bits for suit

# Negative Integers

- Base  $n$  doesn't do negative numbers (if  $n \geq 0$ )
- In written notation, we denote negative numbers with –
- You can designate a bit to be a sign bit
  - This is *sign-magnitude* representation
- In 4 bits,  $0100 = 4$ ,  $1100 = -4$ ,  $0110 = 6$ ,  $1110 = -6$
- But  $1000 = -0$  — two zero values!
- Inconvenient for arithmetic

# 1s' Complement

- Idea: to find  $-x$ , negate the bits in  $x$
- Still two zeroes
- Inconvenient for arithmetic

|     |   |     |    |
|-----|---|-----|----|
| 000 | 0 | 100 | -3 |
| 001 | 1 | 101 | -2 |
| 010 | 2 | 110 | -1 |
| 011 | 3 | 111 | -0 |

# 2's Complement

- Idea: Most significant bit is negative
  - In 3 bits,  $101_2 = -4 + 0 + 1 = -3$
- To find  $-x$ , negate the bits in  $x$  and add 1
- Only one zero, but extra minimum value
- $+$ ,  $-$ ,  $\times$  work without changes
- Used on most computers

|     |   |     |    |
|-----|---|-----|----|
| 000 | 0 | 100 | -4 |
| 001 | 1 | 101 | -3 |
| 010 | 2 | 110 | -2 |
| 011 | 3 | 111 | -1 |

# Range of 2's Complement

- Given  $k$  bits, you can represent  $2^k$  values
- If they are natural numbers, that is  $0 - (2^k - 1)$
- What is the range for 2's complement integers?
  - $x = -d_{k-1}2^{k-1} + d_{k-2}2^{k-2} + \dots d_12^1 + d_02^0$
  - $x$  is negative iff  $d_{k-1} = 1$
- Minimum value is  $1000\dots 0$ , i.e.,  $-(2^{k-1})$
- Maximum value is  $0111\dots 1$ , i.e.,  $2^{k-1} - 1$

# iClicker Quiz

# iClicker Quiz 1

- What is the largest unsigned integer in  $k$  bits?
  - A.  $2^k$
  - B.  $2^k - 1$
  - C.  $2^{k-1}$
  - D.  $2^{k-1} - 1$
  - E.  $k$

# iClicker Quiz 2

- What is the smallest unsigned integer in  $k$  bits?
  - A.  $2^k$
  - B.  $2^k - 1$
  - C.  $-2^{k-1}$
  - D.  $-2^k + 1$
  - E. 0



# iClicker Quiz 3

- What is the largest 2's complement integer in  $k$  bits?
  - A.  $2^k$
  - B.  $2^k - 1$
  - C.  $2^{k-1}$
  - D.  $2^{k-1} - 1$
  - E.  $k$

# iClicker Quiz 4

- What is the smallest 2's complement integer in  $k$  bits?

A.  $2^k$

B.  $2^k - 1$

C.  $-2^{k-1}$

D.  $-2^k + 1$

E. 0

# 2's Complement Addition & Subtraction

Addition: same as unsigned binary addition

Subtraction: invert subtrahend and add

$$\begin{array}{r} -7 \quad 1001 \\ + \quad 5 \quad 0101 \\ \hline -2 \quad 1110 \end{array}$$

$$\begin{array}{r} -5 \quad 1011 \\ + -2 \quad 1110 \\ \hline -7 \quad 11001 \end{array}$$

$$4 - 2 = 4 + (-2)$$

$$\begin{array}{r} \quad 4 \quad 0100 \\ + -2 \quad 1110 \\ \hline \quad 2 \quad 10010 \end{array}$$

Ignore carries

# 2's Complement Overflow

- Addition results can be too large to fit in  $k-1$  bits
  - Since we are using regular addition, we overflow into the sign bit

$$\begin{array}{r} 7 \quad 0111 \\ + \quad 5 \quad 0101 \\ \hline -4 \quad 1100 \end{array}$$

$$\begin{array}{r} -7 \quad 1001 \\ + \quad -7 \quad 1001 \\ \hline 2 \quad 10010 \end{array}$$

- Detecting overflow:
  - positive + positive  $\Rightarrow$  negative
  - negative + negative  $\Rightarrow$  positive

**Does not work in C!**


# Aside: Modular Arithmetic

- In modular arithmetic, we talk about congruence:
  - $a \equiv b \pmod{m}$
  - This means  $\exists d : a = b + md$
- For example:  $9 + 9 \equiv 2 \pmod{16}$ , because  $18 = 2 + 16 \times 1$
- In particular,  $-1 \equiv 15 \pmod{16}$ , because  $-1 = 15 - 16$

# Aside: Modular Arithmetic

- Modular arithmetic operators give congruent values in the range  $0 \dots m-1$ 
  - $a + b \equiv c \pmod{m}$ , where  $0 \leq c < m$
- Modular arithmetic has nearly all the algebraic properties we expect for numbers
  - $+$  and  $\times$  are commutative and associative
  - $(a + b) \times c = a \times c + b \times c$
  - $a - b = a + (-b)$

# Modular Numbers



- $m+n$  — start at  $m$ , move  $n$  spaces clockwise
- $m-n$  — start at  $m$ , move  $n$  spaces counter-clockwise
- Moving  $n$  spaces is the same as moving  $8-n$  spaces the other way

# Multiplication

```
      1011  (multiplicand)
×      101  (multiplier)
-----
      1011  (multiplicand × 1)
shift left 0
      0000  (multiplicand × 0)
shift left 1
      1011  (multiplicand × 1)
shift left 2
-----
  110111  (product)
```



# Algorithm

1.  $\text{result} \leftarrow 0$
  2. If LSB of multiplier = 1, add multiplicand to result
  3. Shift multiplicand left 1 bit (fill LSB with 0)
  4. Shift multiplier right 1 bit (fill MSB with 0)
  5. If multiplier  $> 0$ , go to 2
- We only need to know add and shift to multiply

# Divison

|      |      |        |                   |
|------|------|--------|-------------------|
|      |      | 101    |                   |
|      |      | ÷----- |                   |
| 15÷3 | 0011 | 1111   |                   |
|      | 1100 | 0100   | -(divisor << 2)×1 |
|      |      | -----  |                   |
|      |      | 10011  |                   |
|      | 0110 | 0000   | -(divisor << 1)×0 |
|      |      | -----  |                   |
|      |      | 0011   |                   |
|      | 0011 | 1101   | -(divisor << 0)×1 |
|      |      | -----  |                   |
|      |      | 10000  | remainder         |

# (Un)signed Ints in C

- Unsigned values in C
  - type: `unsigned int i = 10;`
  - cast: `i = (unsigned) j;`
- Casting leaves bits unchanged
  - `(unsigned) -1  $\Rightarrow 2^k - 1$`

# Sign Extension

| Signed Integer | 4 bit | 8 bit    | 16 bit           |
|----------------|-------|----------|------------------|
| <b>+1</b>      | 0001  | 00000001 | 0000000000000001 |
| <b>-1</b>      | 1111  | 11111111 | 1111111111111111 |

# Bit Shifting

- C operators `<<` and `>>` shift an integer by some number of places
- `n << c` returns `n` shifted `c` places left
  - $n \ll c = n \times 2^c$
- `n >> c` returns `n` shifted `c` places right
  - $n \gg c = n \div 2^c$
- What about sign?
  - `>>`, `/`, and `%` operate differently for signed and unsigned ints

# Base -2

- Idea: even places are positive, odd are negative
- Extremely unbalanced
- No sign extension
- Arithmetic is complicated—even negation
- (Probably) never used anywhere

|     |    |     |   |
|-----|----|-----|---|
| 000 | 0  | 100 | 4 |
| 001 | 1  | 101 | 5 |
| 010 | -2 | 110 | 2 |
| 011 | -1 | 111 | 3 |

# Floating Point

# Scientific Notation

- Another way of writing numbers
  - $1,250,000 = 1.25 \times 10^6$
  - Can be more compact
  - Distinguish magnitude from precision
- Can be used with base  $n$  for any  $n$ 
  - $110001.1_2 \approx 1.1_2 \times 2^5$ ;  $0.00101_2 = 1.01_2 \times 2^{-3}$
- How might we represent this in computers?



# Fixed-Point Numbers

- Choose an exponent  $e$  and represent  $m \times b^e$  as an integer
  - For example,  $\$10.50 = \$1050 \times 10^{-2}$
- Use integer operations to add and subtract
- To multiply, use integer multiplication and then multiply by  $b^e$ 
  - $(45 \times 10^{-2})(1000 \times 10^{-2}) = (45 \times 1000 \times 10^{-2}) \times 10^{-2}$   
 $= 450 \times 10^{-2}$

# Floating-Point Numbers

- Idea: Represent numbers as  $m \times b^e$
- Represent  $m$  and  $e$  as (fixed-width) integers
  - Base  $b$  is fixed (usually 2, sometimes 10)
- Allows for a much wider range of values
- Variable precision
  - Values closer to 0 are more precise — why?

# IEEE 754 Floating-Point

- The most commonly-used representation for floating-point
- Designed by experts in numerics
  - Has lots of features that surprise non-experts
- Defines 3 standard precisions for binary FP:
  - single (32 bit), double (64 bit), extended (80 bit)
- Some processors also support half (16 bit) or quad (128 bit)

# Normal Values

- All IEEE FP numbers have three parts
  - sign  $s$ , exponent  $e$ , significand  $m$
  - $s$  is 1 bit;  $e$  and  $m$  depend on the precision
- Conceptually:  $(-1)^s \times m \times 2^{e-\text{bias}(E)}$ , where  $1 \leq m < 2$ 
  - $\text{bias}(E) = 2^{E-1}-1$ , where  $E$  is the number of bits in  $e$
  - Yet another way to represent signed numbers
  - $e \neq 0$  and  $e \neq 2^E-1$  (e.g., all ones)

# Other Values

- Zero and “denormal” values — less than any normal value
  - Conceptually:  $(-1)^s \times m \times 2^{1-\text{bias}(E)}$ , where  $0 \leq m < 1$
  - $\pm 0$  if  $m = 0$  and  $e = 0$
  - Denormal if  $m \neq 0$  and  $e = 0$
- Non-finite values
  - $\pm\infty$  if  $m = 0$  and  $e = 2^E - 1$  (all ones)
  - NaN if  $m \neq 0$  and  $e = 2^E - 1$

# Conversion to FP

- Single precision:  $E = 8$ ,  $M = 23$  (mantissa bits)
- Rewrite 5.625 in binary:  $101.101_2$
- Rewrite in scientific notation:  $1.01101_2 \times 2^2$
- Solve  $e - \text{bias}(E) = 2$ 
  - In single precision,  $\text{bias}(E) = 127$ , so  $e = 129 = 10000001_2$
- Drop the integer part of  $m$  (we can deduce it using  $e$ ) and extend to  $M$  bits:  $01101000000000000000000_2$
- Thus: 0 10000001 011010000000000000000000

# Question 1

If 1011 is a 4-bit, 2's complement integer, what is its value?

A. 11

B. 0

C. -3

D. -4

E. -5

# Question 2

IEEE double-precision floating point has how many bits?

- A. 1
- B. 16
- C. 23
- D. 32
- E. 64



# Question 3

What is  $1.101 \times 2^1$  in decimal?

A. 1.625

B. 3.25

C. 6.5

D. 13

E. -3.75

# Question 4

What was the hardest part of PA1?

- A. Memory management
- B. Input/output
- C. Using the command-line tools
- D. Accessing the iLab
- E. Time management